

Problem Set - 2

Arijit Das, Roll - 713

2026-02-02

Linear Regression

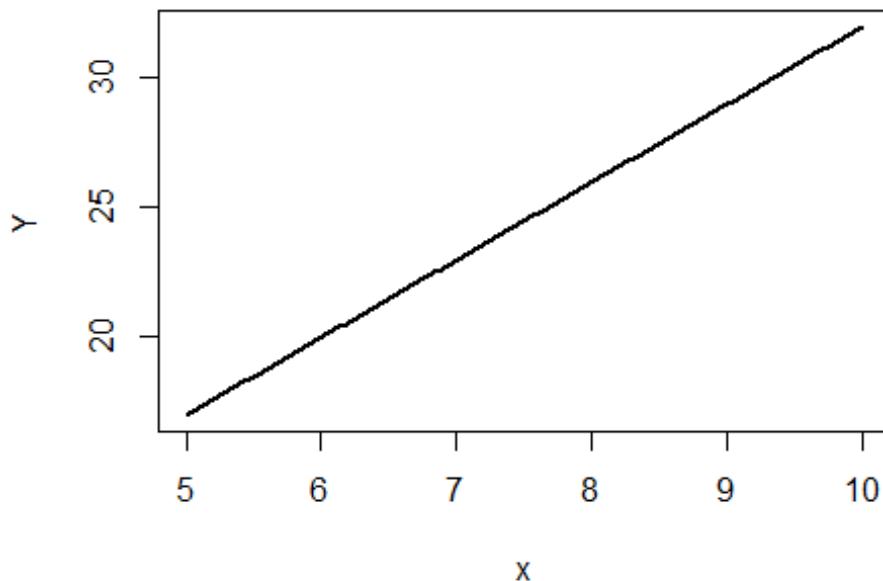
1. Problem to demonstrate that the population regression line is fixed, but least square regression line varies.

Suppose the population regression line is given by $Y = 2 + 3x$, while the data comes from the model $y = 2 + 3x + \epsilon$.

Step 1: For x in the range [5, 10] graph the population regression line.

```
rm(list=ls())
x=seq(5,10,length.out=200)
y=2+3*x
plot(x,y,type='l',lwd=2,main="PRF of Y=2+3x",xlab = "x",ylab = "Y")
```

PRF of $Y=2+3x$



Step 2: Generate $x_i(i = 1, 2, \dots, n)$ from $\text{Uniform}(5, 10)$ and $\epsilon_i(i = 1, 2, \dots, n)$ from $N(0, 4^2)$. Hence, compute y_1, y_2, \dots, y_n .

```
set.seed(123)
n=50
xi=runif(n,5,10)
```

```

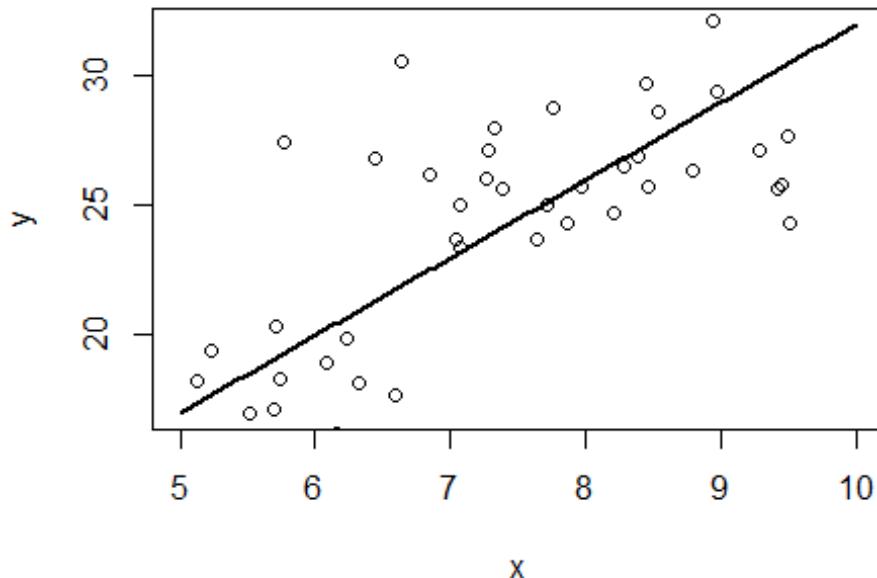
ei=rnorm(n,0,4)
yi=2+3*xi+ei
head(data.frame(xi,ei,yi))

##           xi          ei        yi
## 1 6.437888 -6.7467732 14.56689
## 2 8.941526  3.3511482 32.17573
## 3 7.044885  0.6134925 23.74815
## 4 9.415087 -4.5525477 25.69271
## 5 9.702336  5.0152597 36.12227
## 6 5.227782  1.7058569 19.38920

plot(x,y,type='l',lwd=2,main="Plot of the PRF Y=2+3x and the Sample points")
points(xi,yi)

```

Plot of the PRF $Y=2+3x$ and the Sample points



Step 3: On the basis of the data (x_i, y_i) ($i = 1, 2, \dots, n$) generated in Step 2, report the least squares regression line.

```

model=lm(yi~xi)
summary(model)

##
## Call:
## lm(formula = yi ~ xi)
##
## Residuals:
##      Min       1Q   Median       3Q      Max 
## -10.0000 -2.0000  0.0000  2.0000 10.0000
## 
```

```

## -9.0231 -2.2314 -0.2627  2.1970  8.7445
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)
## (Intercept) -0.09639   2.82610  -0.034    0.973
## xi          3.30540   0.36519   9.051 5.96e-12 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 3.761 on 48 degrees of freedom
## Multiple R-squared:  0.6306, Adjusted R-squared:  0.6229
## F-statistic: 81.93 on 1 and 48 DF,  p-value: 5.962e-12

b0=coef(model)[1];b0

## (Intercept)
## -0.09638929

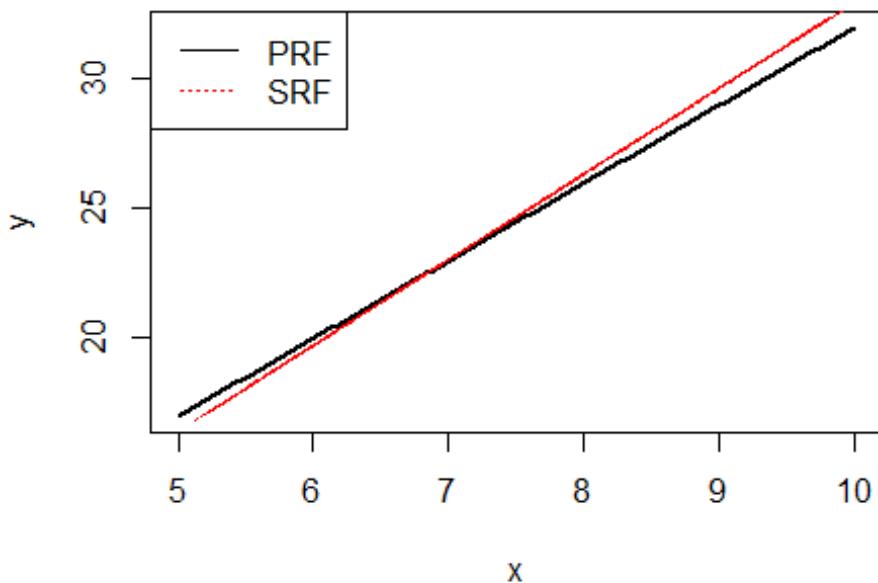
b1=coef(model)[2];b1

##      xi
## 3.305396

y.hat=b0+b1*xi
plot(x,y,type='l',lwd=2,main="PRF and SRF")
lines(xi,y.hat,col="red",lty="dotted")
legend("topleft",col=c("black","red"),legend=c("PRF","SRF"),
       lty=c("solid","dotted"))

```

PRF and SRF



Step 4: Repeat steps 2-3 five times. Graph the 5 least squares regression lines over the population regression line obtained in Step 1.

```
rm(list=ls())
set.seed(123)
n = 50
x_grid = seq(5,10,length.out=200)
y_grid = 2 + 3*x_grid

coefs = matrix(0, 5, 2)

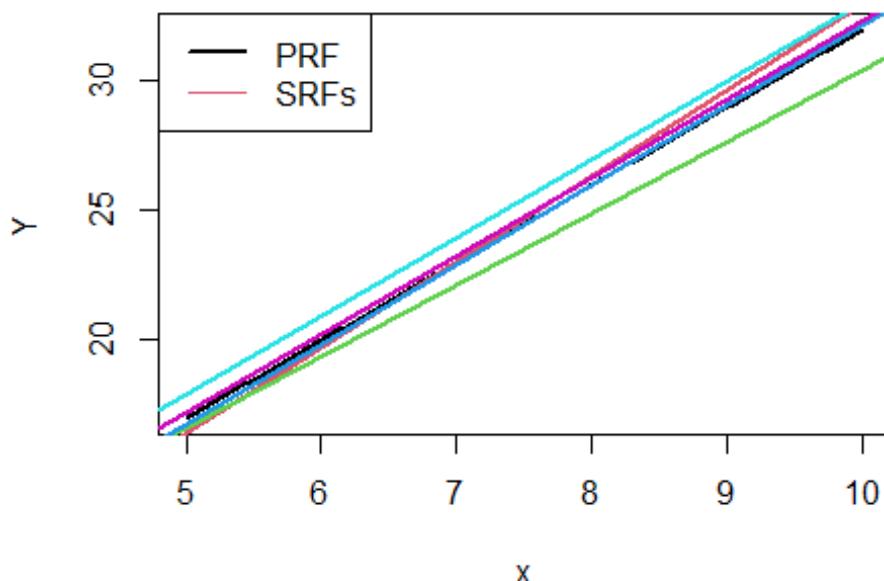
plot(x_grid, y_grid, type='l', lwd=2, col="black",main="PRF and SRFs",
xlab="x", ylab="Y")

for(i in 1:5){

  x_sim = runif(n,5,10)
  e = rnorm(n,0,4)
  y_sim = 2 + 3*x_sim + e
  model = lm(y_sim ~ x_sim)
  coefs[i,] = coef(model)
  abline(model, col=i+1, lwd=2)
}

legend("topleft",legend=c("PRF","SRFs"),col=c("black",2),lty=c(1,1),lwd=c(2,1))
```

PRF and SRFs



```

coefs

##          [,1]      [,2]
## [1,] -0.09638929 3.305396
## [2,]  2.79218839 2.761042
## [3,]  1.39299737 3.073267
## [4,]  2.82308856 3.023608
## [5,]  2.03250638 3.028097

```

The PRF remains same. But the SRF is not fixed

2 Problem to demonstrate that $\hat{\beta}^0$ and $\hat{\beta}^$ minimises RSS*

Step 1: Generate x_i from $Uniform(5, 10)$ and mean centre the values. Generate ϵ_i from $N(0, 1)$. Calculate $y_i = 2 + 3x_i + \epsilon_i$, $i = 1, 2, \dots, n$. Take $n=50$ and seed=123

```

rm(list=ls())
n=50
set.seed(123)
x=runif(n,5,10)
xm=x-mean(x)
e=rnorm(n)
y=2+3*xm+e

```

Step 2: Now imagine that you only have the data on (x_i, y_i) , $i = 1, 2, \dots, n$, without knowing the mechanism that was used to generate the data in step 1. Assuming a linear regression of the type $y_i = \beta_0 + \beta x_i + \epsilon_i$, and based on these data (x_i, y_i) , $i = 1, 2, \dots, n$, obtain the least squares estimates of β_0 and β .

```

fit=lm(y~xm)
summary(fit)

##
## Call:
## lm(formula = y ~ xm)
##
## Residuals:
##       Min     1Q   Median     3Q    Max 
## -2.25578 -0.55786 -0.06567  0.54926  2.18613 
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)    
## (Intercept)  2.0562    0.1330  15.46   <2e-16 ***
## xm          3.0764    0.0913  33.70   <2e-16 ***
## ---        
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1 
##
## Residual standard error: 0.9404 on 48 degrees of freedom
## Multiple R-squared:  0.9594, Adjusted R-squared:  0.9586 
## F-statistic: 1135 on 1 and 48 DF,  p-value: < 2.2e-16

```

```

beta0=coef(fit)[1];beta0
## (Intercept)
##      2.056189

beta=coef(fit)[2];beta
##      xm
##  3.076349

```

Step 3: Take a large number of grid values of (β_0, β) that also include the least squares estimates obtained from step 2. Compute the RSS for each parametric choice of (β_0, β) , where $RSS = (y_1 - \beta_0 - \beta x_1)^2 + (y_2 - \beta_0 - \beta x_2)^2 + \dots + (y_n - \beta_0 - \beta x_n)^2$. Find out for which combination of (β_0, β) , RSS is minimum.

```

beta0.grid=seq(-1,1,length.out=101)+beta0
beta.grid=seq(-1,1,length.out=101)+beta

RSS=c()
for(i in 1:length(beta.grid))
{
  RSS[i]=sum((y-beta0.grid[i]-beta.grid[i]*xm)^2)
}
RSS

## [1] 198.53732 192.35609 186.29972 180.36823 174.56162 168.87987 163.32301
## [8] 157.89101 152.58389 147.40164 142.34426 137.41176 132.60413 127.92138
## [15] 123.36350 118.93049 114.62236 110.43910 106.38071 102.44719 98.63855
## [22] 94.95479 91.39589 87.96187 84.65273 81.46845 78.40905 75.47453
## [29] 72.66487 69.98009 67.42019 64.98516 62.67500 60.48971 58.42930
## [36] 56.49376 54.68310 52.99730 51.43639 50.00034 48.68917 47.50287
## [43] 46.44145 45.50490 44.69322 44.00641 43.44448 43.00743 42.69524
## [50] 42.50793 42.44550 42.50793 42.69524 43.00743 43.44448 44.00641
## [57] 44.69322 45.50490 46.44145 47.50287 48.68917 50.00034 51.43639
## [64] 52.99730 54.68310 56.49376 58.42930 60.48971 62.67500 64.98516
## [71] 67.42019 69.98009 72.66487 75.47453 78.40905 81.46845 84.65273
## [78] 87.96187 91.39589 94.95479 98.63855 102.44719 106.38071 110.43910
## [85] 114.62236 118.93049 123.36350 127.92138 132.60413 137.41176 142.34426
## [92] 147.40164 152.58389 157.89101 163.32301 168.87987 174.56162 180.36823
## [99] 186.29972 192.35609 198.53732

df=data.frame(beta0.grid,beta.grid,RSS)
df

##   beta0.grid beta.grid      RSS
## 1     1.056189  2.076349 198.53732
## 2     1.076189  2.096349 192.35609
## 3     1.096189  2.116349 186.29972
## 4     1.116189  2.136349 180.36823
## 5     1.136189  2.156349 174.56162
## 6     1.156189  2.176349 168.87987
## 7     1.176189  2.196349 163.32301

```

## 8	1.196189	2.216349	157.89101
## 9	1.216189	2.236349	152.58389
## 10	1.236189	2.256349	147.40164
## 11	1.256189	2.276349	142.34426
## 12	1.276189	2.296349	137.41176
## 13	1.296189	2.316349	132.60413
## 14	1.316189	2.336349	127.92138
## 15	1.336189	2.356349	123.36350
## 16	1.356189	2.376349	118.93049
## 17	1.376189	2.396349	114.62236
## 18	1.396189	2.416349	110.43910
## 19	1.416189	2.436349	106.38071
## 20	1.436189	2.456349	102.44719
## 21	1.456189	2.476349	98.63855
## 22	1.476189	2.496349	94.95479
## 23	1.496189	2.516349	91.39589
## 24	1.516189	2.536349	87.96187
## 25	1.536189	2.556349	84.65273
## 26	1.556189	2.576349	81.46845
## 27	1.576189	2.596349	78.40905
## 28	1.596189	2.616349	75.47453
## 29	1.616189	2.636349	72.66487
## 30	1.636189	2.656349	69.98009
## 31	1.656189	2.676349	67.42019
## 32	1.676189	2.696349	64.98516
## 33	1.696189	2.716349	62.67500
## 34	1.716189	2.736349	60.48971
## 35	1.736189	2.756349	58.42930
## 36	1.756189	2.776349	56.49376
## 37	1.776189	2.796349	54.68310
## 38	1.796189	2.816349	52.99730
## 39	1.816189	2.836349	51.43639
## 40	1.836189	2.856349	50.00034
## 41	1.856189	2.876349	48.68917
## 42	1.876189	2.896349	47.50287
## 43	1.896189	2.916349	46.44145
## 44	1.916189	2.936349	45.50490
## 45	1.936189	2.956349	44.69322
## 46	1.956189	2.976349	44.00641
## 47	1.976189	2.996349	43.44448
## 48	1.996189	3.016349	43.00743
## 49	2.016189	3.036349	42.69524
## 50	2.036189	3.056349	42.50793
## 51	2.056189	3.076349	42.44550
## 52	2.076189	3.096349	42.50793
## 53	2.096189	3.116349	42.69524
## 54	2.116189	3.136349	43.00743
## 55	2.136189	3.156349	43.44448
## 56	2.156189	3.176349	44.00641
## 57	2.176189	3.196349	44.69322

```

## 58 2.196189 3.216349 45.50490
## 59 2.216189 3.236349 46.44145
## 60 2.236189 3.256349 47.50287
## 61 2.256189 3.276349 48.68917
## 62 2.276189 3.296349 50.00034
## 63 2.296189 3.316349 51.43639
## 64 2.316189 3.336349 52.99730
## 65 2.336189 3.356349 54.68310
## 66 2.356189 3.376349 56.49376
## 67 2.376189 3.396349 58.42930
## 68 2.396189 3.416349 60.48971
## 69 2.416189 3.436349 62.67500
## 70 2.436189 3.456349 64.98516
## 71 2.456189 3.476349 67.42019
## 72 2.476189 3.496349 69.98009
## 73 2.496189 3.516349 72.66487
## 74 2.516189 3.536349 75.47453
## 75 2.536189 3.556349 78.40905
## 76 2.556189 3.576349 81.46845
## 77 2.576189 3.596349 84.65273
## 78 2.596189 3.616349 87.96187
## 79 2.616189 3.636349 91.39589
## 80 2.636189 3.656349 94.95479
## 81 2.656189 3.676349 98.63855
## 82 2.676189 3.696349 102.44719
## 83 2.696189 3.716349 106.38071
## 84 2.716189 3.736349 110.43910
## 85 2.736189 3.756349 114.62236
## 86 2.756189 3.776349 118.93049
## 87 2.776189 3.796349 123.36350
## 88 2.796189 3.816349 127.92138
## 89 2.816189 3.836349 132.60413
## 90 2.836189 3.856349 137.41176
## 91 2.856189 3.876349 142.34426
## 92 2.876189 3.896349 147.40164
## 93 2.896189 3.916349 152.58389
## 94 2.916189 3.936349 157.89101
## 95 2.936189 3.956349 163.32301
## 96 2.956189 3.976349 168.87987
## 97 2.976189 3.996349 174.56162
## 98 2.996189 4.016349 180.36823
## 99 3.016189 4.036349 186.29972
## 100 3.036189 4.056349 192.35609
## 101 3.056189 4.076349 198.53732

df[which.min(RSS),]

##      beta0.grid beta.grid      RSS
## 51    2.056189  3.076349 42.4455

```

3 Problem to demonstrate that least square estimators are unbiased

Step 1: Generate $x_i(i = 1, 2, \dots, n)$ from $\text{Uniform}(0, 1)$, $\epsilon_i(i = 1, 2, \dots, n)$ from $N(0, 1)$ and hence generate y using $y_i = \beta_0 + \beta x_i + \epsilon_i$. (Take $\beta_0 = 2$, $\beta = 3$).

```
rm(list=ls())
n=50
set.seed(123)
x=runif(n,0,1)
e=rnorm(n,0,1)
beta_not=2
beta=3
y=beta_not+beta*x+e
```

Step 2: On the basis of the data $(x_i, y_i)(i = 1, 2, \dots, n)$ generated in Step 1, obtain the least square estimates of β_0 and β .

```
model=lm(y~x)
b0=coef(model)[1];b0
## (Intercept)
##     1.857647

b1=coef(model)[2];b1
##          x
## 3.381745
```

Repeat Steps 1-2, R = 1000 times. In each simulation obtain $\hat{\beta}_0$ and $\hat{\beta}$. Finally, the least-square estimates will be given by the average of these estimated values. Compare these with the true β_0 and β and comment.

Take $n = 50$ and $\text{seed}=123$.

```
R = 1000
b0_store = c()
b1_store = c()

for(i in 1:R){
  x = runif(n,0,1)
  e = rnorm(n,0,1)
  y = beta_not + beta*x + e
  model = lm(y~x)
  b0_store[i] = coef(model)[1]
  b1_store[i] = coef(model)[2]
}

mean_b0 = mean(b0_store)
mean_b1 = mean(b1_store)
mean_b0
```

```

## [1] 2.013052
mean_b1
## [1] 2.981907

```

4 Comparing several simple linear regressions

Attach “Boston” data from MASS library in R. Select median value of owneroccupied homes, as the response and per capita crime rate, nitrogen oxides concentration, proportion of blacks and percentage of lower status of the population as predictors.

(a) Selecting the predictors one by one, run four separate linear regressions to the data. Present the output in a single table.

(b) Which model gives the best fit?

(c) Compare the coefficients of the predictors from each model and comment on the usefulness of the predictors.

```

rm(list=ls())
library(MASS)
head(Boston)

##      crim zn indus chas   nox      rm    age     dis rad tax ptratio black lstat
## 1 0.00632 18  2.31    0 0.538 6.575 65.2 4.0900    1 296 15.3 396.90 4.98
## 2 0.02731  0  7.07    0 0.469 6.421 78.9 4.9671    2 242 17.8 396.90 9.14
## 3 0.02729  0  7.07    0 0.469 7.185 61.1 4.9671    2 242 17.8 392.83 4.03
## 4 0.03237  0  2.18    0 0.458 6.998 45.8 6.0622    3 222 18.7 394.63 2.94
## 5 0.06905  0  2.18    0 0.458 7.147 54.2 6.0622    3 222 18.7 396.90 5.33
## 6 0.02985  0  2.18    0 0.458 6.430 58.7 6.0622    3 222 18.7 394.12 5.21
##      medv
## 1 24.0
## 2 21.6
## 3 34.7
## 4 33.4
## 5 36.2
## 6 28.7

attach(Boston)
model1=lm(medv~crim)
model2=lm(medv~nox)
model3=lm(medv~black)
model4=lm(medv~lstat)
library(stargazer)

stargazer(model1,model2,model3,model4,type="html",out="f1.html")

```

	<i>Dependent variable:</i> medv			
	(1)	(2)	(3)	(4)
crim	-0.415 *** (0.044)			
nox		-33.916 *** (3.196)		
black			0.034 *** (0.004)	
lstat				-0.950 *** (0.039)
Constant	24.033 *** (0.409)	41.346 *** (1.811)	10.551 *** (1.557)	34.554 *** (0.563)
Observations	506	506	506	506
R ²	0.151	0.183	0.111	0.544
Adjusted R ²	0.149	0.181	0.109	0.543
Residual Std. Error (df = 504)	8.484	8.323	8.679	6.216
F Statistic (df = 1; 504)	89.486 ***	112.591 ***	63.054 ***	601.618 ***

Note: *p<0.1; **p<0.05; ***p<0.01